

Heaps and Priority Queues

CS 2860: Algorithms and Complexity

Magnus Wahlström and Gregory Gutin

February 17, 2019

Heaps and Priority Queues

The Priority Queue ADT

- ▶ A collection of items, each with a **priority**
- ▶ Exists in two variants:
 1. **Min-priority queue**:
 - ▶ **insert(item)**: Add `item` to queue
 - ▶ **minItem()**: Return the item with the **smallest priority**
 - ▶ **deleteMin()**: Remove and return (pop) the item with the **smallest priority** from the queue
 2. **Max-priority queue**:
 - ▶ **insert(item)**: Add `item` to queue
 - ▶ **maxItem()**: Return the item with the **largest** priority
 - ▶ **deleteMax()**: Remove and return (pop) the item with the **largest** priority from the queue
- ▶ **No support** for efficient random access, ordered iteration, etc.
- ▶ Optional advanced (but useful) operation:
 - ▶ **changePriority(item, priority)**: Update priority of item already in the priority queue

Priority Queues: Examples

- ▶ For max-PQ: **job pool**:
 - ▶ Insert jobs to be executed with some “importance” (priority) parameter
 - ▶ Pull out **most important** job to perform next
- ▶ For min-PQ: **scheduling**:
 - ▶ Insert events with **triggering times** (e.g., at 12:05:00, launch task T)
 - ▶ Look at **most imminent** event to execute
- ▶ Common tool in algorithms (A^* search, Dijkstra, Prim):
- ▶ Common pattern:
 1. Start with initial state S
 2. Discover some new states from S , add to PQ
 3. Select the **best** (min-cost) discovered state S' as our next state
 4. Discover **new** states from S' , add to PQ
 5. Repeat until done

Questions

1. Assume we are using a **max-PQ** with objects with an **item.priority** field storing the priority value.

How does the PQ operate in the following cases?

- 1.1 If every new item has priority **larger** than all previous ones?
(E.g., priority is a ticking “clock”.)

- ▶ **pq.deleteMax()** acts like **stack.pop()** – newest first

- 1.2 If every new item has priority **smaller** than all previous ones?

- ▶ **pq.deleteMax()** acts like **queue.dequeue()** – oldest first

2. Which data structures that we already know could provide good/bad implementations of a Priority Queue?

Some options:

- ▶ Same as a **Queue** (linked list/array): always possible, but bad worst-case performance
- ▶ Self-balancing BST, e.g., **AVL tree** or **Red/black tree**: good worst-case bounds, very complex code.
- ▶ Will see **binary heaps** data structure

Implementation 0: As Queue

- ▶ It's **possible** to implement PQ via a linked list:
 - ▶ Linked list contains items sorted by priority
 - ▶ `minItem()` and `deleteMin()` work in $\Theta(1)$ time (just look at / delete the head of the list)
 - ▶ `insert(item)`: insert in correct position (up to $\Theta(n)$ work to insert into middle position)
- ▶ Array-based variant: **circular buffer** (very similar profile)
- ▶ These **may** be efficient if items arrive **mostly in sequence**
- ▶ ...but we want **worst-case guarantees**, which this doesn't give

Implementation 1: Trees

- ▶ We can easily support a Priority Queue with a balanced binary search tree:
 - ▶ insert: Exists, time $O(\log n)$
 - ▶ minValue: Exists, time $O(\log n)$
 - ▶ deleteMin: Exists via min+delete, time $O(\log n)$
- ▶ Even the operation `changePriority(item, priority)`:
 - ▶ Delete old item, insert with new priority
 - ▶ $O(\log n)$ total time

Implementation 2: Heaps

- ▶ Canonical implementation of priority queues
- ▶ Operations: insert, deleteMin in $O(\log n)$ time, minVal in $O(1)$ time
- ▶ Can bootstrap: $O(n)$ -time initialisation with n items

Why heaps?

Running times of tree and heap implementations nearly **identical** – why heaps?

- ▶ Heaps live in **arrays** – trees are linked structures (use nodes)
 - ▶ Heaps have better memory usage
 - ▶ Heaps are slightly faster in practice (tighter code)
- ▶ Simpler code
- ▶ Offer interesting **extensions** (beyond scope of this course)

We will see...

- ▶ [Now] Heaps: What is a heap? How is it represented in memory?
- ▶ [Now] Heaps: Implementing `insert`, `deleteMin`
- ▶ [Later] $O(n)$ -time `makeHeap`
- ▶ [Later] The heapsort algorithm

Binary Heaps

Heaps

1. A heap is a tree with a special property
2. A **binary heap** is a heap represented as an array
 - ▶ **Logically**: Works like a tree
 - ▶ **Physically**: Lives in memory like an array (no pointers)
3. “Heap tree” property:
 - ▶ Every node in the tree is smaller than all its children
 - ▶ By induction: Smallest item in root
 - ▶ **Not** a binary **search** tree!

Binary heaps: Tree/array representation

- ▶ To map tree into array, each node gets a **number** (levelorder)
 - ▶ Root is number 0
 - ▶ The children of node i are number $2i + 1$ and $2i + 2$
 - ▶ Can access data for node i as `array[i]` (no links needed)
- ▶ A binary heap is always **nearly complete**
 - ▶ All nodes $0, 1, \dots, n - 1$ exist (no “holes” in array)
 - ▶ Implies: All nodes have exactly two children, except possibly at last and second-last level

Main heap property

Heap property

Every node in the tree contains a smaller item than both its children

- ▶ By induction:
 - ▶ **Smallest** item always in the root
 - ▶ **Largest** item can be **anywhere** on leaf level (among $n/2$ options)
- ▶ More relaxed than a search tree – easier to maintain

Heap operations


Presented algorithms

1. Primitive (underlying) heap operations:
 - ▶ `siftUp(index)` (“swim”), `siftDown(index)` (“sink”)
2. Fundamental Priority Queue operations:
 - ▶ `min`, `insert`, `deleteMin`
3. `makeHeap` operation: Fast initialisation
4. `heapsort` sorting algorithm

Primitive 1: siftUp

- ▶ Case: Tree is **almost** a heap, except **one item** (a leaf) that is in the wrong place
- ▶ Need to “sift up”: **bubble** item upwards until it finds its place
- ▶ Code:
 1. While `array[i] < array[parent(i)]`:
 - 1.1 Swap contents `array[i]` and `array[parent(i)]`
 - 1.2 Let `i=parent(i)` (break if trying to find parent of node 1)

Sift up: Visualisation




Sift up: Analysis

- ▶ Question 1: Why does it work? (Why is the result a heap?)
 - ▶ Assumed: Only problem is that node i is out of place
 - ▶ In swap, only shifted i and $\text{parent}(i)$ around
 - ▶ Easy to check that i and $\text{parent}(i)$ are both smaller than both their children after the swap
- ▶ Question 2: What is its running time?
 - ▶ Number of steps \leq height of tree
 - ▶ Tree balanced $\Rightarrow O(\log n)$ time

Primitive 2: Sift down

- ▶ Case: Tree is (again) almost a heap, except one item (maybe the root) is **too large**
- ▶ Need to “sift” or “bubble” item **downwards** until it finds its place
- ▶ Code `siftDown(i)`:
 1. Let `smallest` be smallest node of `i`, `leftChild(i)`, `rightChild(i)`
 2. If `smallest` \neq `i`:
 - 2.1 Swap contents `array[i]` and `array[smallest]`
 - 2.2 Call `siftDown(smallest)` recursively

Sift down: Visualisation



Sift down: Analysis (brief)

1. Correctness: Similar to `siftUp` (somewhat longer arguments)
2. Running time: $O(\text{tree height}) = O(\log n)$.

Priority Queue operations

- ▶ Operation `min`:
 1. Return `array[0]`
- ▶ Operation `insert(item)`:
 1. Add `item` to the end of array
 2. Call `siftUp(n)` to put it in its place
- ▶ Operation `deleteMin`:
 1. Swap root with last item of tree
 2. “Forget” last item (let $n=n-1$)
 3. Call `siftDown(0)` to put other item in place

Priority Queue: Summary

- ▶ Priority Queue: Useful ADT for situations where we want to frequently access the “most urgent” / “best-looking” item
- ▶ Well-implemented (canonically) by the **heap** structure
- ▶ Advantages:
 - ▶ Small memory usage (no extra space for “links” / pointers or node objects)
 - ▶ Simple code, fast operations
- ▶ However: Need to understand heap structure (“tree as array”)
- ▶ Extensions exist (binomial heap, Fibonacci heap) – technically, we only saw the **binary heap**