

Minimum Spanning Trees

CS 2860: Algorithms and Complexity

Magnus Wahlström and Gregory Gutin

February 24, 2018

Minimum Spanning Trees

Minimum spanning trees

- ▶ Input: Graph $G = (V, E)$, edge weights/costs $w(e)$
- ▶ **Spanning tree**: A tree $T = (V, F)$ with edges from E that connects every vertex of G
- ▶ **Minimum spanning tree**: A spanning tree T of minimum total cost

$$\sum_{e \in E(T)} w(e)$$

Minimum spanning trees

- ▶ Input: Graph $G = (V, E)$, edge weights/costs $w(e)$
- ▶ **Spanning tree**: A tree $T = (V, F)$ with edges from E that connects every vertex of G
- ▶ **Minimum spanning tree**: A spanning tree T of minimum total cost

$$\sum_{e \in E(T)} w(e)$$

- ▶ Historical example: **Electrifying Moravia**
 - ▶ Region of present Czech republic
 - ▶ Build **power lines** to connect **villages** to electric grid
 - ▶ Different connections have **different costs** to build (long/short distance, good/bad terrain)
 - ▶ Borůvka's algorithm, 1926
- ▶ Computer networks, road networks, ...

- ▶ Historical example: **Electrifying Moravia**
 - ▶ Region of present Czech republic
 - ▶ Build **power lines** to connect **villages** to electric grid
 - ▶ Different connections have **different costs** to build (long/short distance, good/bad terrain)
 - ▶ Borůvka's algorithm, 1926
- ▶ Computer networks, road networks, ...

- ▶ Historical example: **Electrifying Moravia**
 - ▶ Region of present Czech republic
 - ▶ Build **power lines** to connect **villages** to electric grid
 - ▶ Different connections have **different costs** to build (long/short distance, good/bad terrain)
 - ▶ Borůvka's algorithm, 1926
- ▶ Computer networks, road networks, ...

Definition

A tree is a connected graph with no cycles

We could also have chosen any of the following (equivalent statements):

- ▶ A tree is a connected graph with $n - 1$ edges
- ▶ A tree is a graph with $n - 1$ edges and no cycles
- ▶ A tree is a graph with a **unique** path between any two vertices
- ▶ A tree is a connected graph such that removing any edge leaves a disconnected graph

Definition

A tree is a connected graph with no cycles

We could also have chosen any of the following (equivalent statements):

- ▶ A tree is a connected graph with $n - 1$ edges
- ▶ A tree is a graph with $n - 1$ edges and no cycles
- ▶ A tree is a graph with a **unique** path between any two vertices
- ▶ A tree is a connected graph such that removing any edge leaves a disconnected graph

Definition

A tree is a connected graph with no cycles

We could also have chosen any of the following (equivalent statements):

- ▶ A tree is a connected graph with $n - 1$ edges
- ▶ A tree is a graph with $n - 1$ edges and no cycles
- ▶ A tree is a graph with a **unique** path between any two vertices
- ▶ A tree is a connected graph such that removing any edge leaves a disconnected graph

Definition

A tree is a connected graph with no cycles

We could also have chosen any of the following (equivalent statements):

- ▶ A tree is a connected graph with $n - 1$ edges
- ▶ A tree is a graph with $n - 1$ edges and no cycles
- ▶ A tree is a graph with a **unique** path between any two vertices
- ▶ A tree is a connected graph such that removing any edge leaves a disconnected graph

Definition

A tree is a connected graph with no cycles

We could also have chosen any of the following (equivalent statements):

- ▶ A tree is a connected graph with $n - 1$ edges
- ▶ A tree is a graph with $n - 1$ edges and no cycles
- ▶ A tree is a graph with a **unique** path between any two vertices
- ▶ A tree is a connected graph such that removing any edge leaves a disconnected graph

Properties of Spanning Trees

Two properties of spanning trees

Let T be a spanning tree in a graph $G = (V, E)$.

► **Cycle** property:

1. **Adding** any edge uv to T creates a **cycle** C
2. **Removing** any edge $u'v'$ from C creates a **new spanning tree**

► **Cut** property:

1. **Removing** any edge uv from T disconnects T into **two parts**
2. **Adding** any edge $u'v'$ between these parts creates a **new spanning tree**

Proof (both cases): The new tree is a cycle-free graph with $n - 1$ edges

Properties of Spanning Trees

Two properties of spanning trees

Let T be a spanning tree in a graph $G = (V, E)$.

► **Cycle** property:

1. **Adding** any edge uv to T creates a **cycle** C
2. **Removing** any edge $u'v'$ from C creates a **new spanning tree**

► **Cut** property:

1. **Removing** any edge uv from T disconnects T into **two parts**
2. **Adding** any edge $u'v'$ between these parts creates a **new spanning tree**

Proof (both cases): The new tree is a cycle-free graph with $n - 1$ edges

Properties of Spanning Trees

Two properties of spanning trees

Let T be a spanning tree in a graph $G = (V, E)$.

► **Cycle** property:

1. **Adding** any edge uv to T creates a **cycle** C
2. **Removing** any edge $u'v'$ from C creates a **new spanning tree**

► **Cut** property:

1. **Removing** any edge uv from T disconnects T into **two parts**
2. **Adding** any edge $u'v'$ between these parts creates a **new spanning tree**

Proof (both cases): The new tree is a cycle-free graph with $n - 1$ edges

Min-cost spanning trees

Min-cost spanning trees: Cycle rule

Let $G = (V, E)$ be a graph with edge weights, T a spanning tree in G . Then T is **min-cost** if and only if the following holds:

- ▶ Let $uv \in E$ be **any** edge not used in T
- ▶ Let C be the cycle created by adding uv to T
- ▶ Then uv is the **most expensive edge** of C
(or tied for most expensive, if there are ties)

Why?

- ▶ Suppose not. Let uv be an edge not used in T , that is not the most expensive in its cycle C in $T + uv$
- ▶ Create a tree T' by adding uv , then deleting the most expensive edge $u'v'$ from C .
- ▶ Then T' is a cheaper spanning tree than T , a contradiction.

Min-cost spanning trees

Min-cost spanning trees: Cycle rule

Let $G = (V, E)$ be a graph with edge weights, T a spanning tree in G . Then T is **min-cost** if and only if the following holds:

- ▶ Let $uv \in E$ be **any** edge not used in T
- ▶ Let C be the cycle created by adding uv to T
- ▶ Then uv is the **most expensive edge** of C (or tied for most expensive, if there are ties)

Why?

- ▶ Suppose not. Let uv be an edge not used in T , that is not the most expensive in its cycle C in $T + uv$
- ▶ Create a tree T' by adding uv , then deleting the most expensive edge $u'v'$ from C .
- ▶ Then T' is a cheaper spanning tree than T , a contradiction.

Min-cost spanning trees

Min-cost spanning trees: Cycle rule

Let $G = (V, E)$ be a graph with edge weights, T a spanning tree in G . Then T is **min-cost** if and only if the following holds:

- ▶ Let $uv \in E$ be **any** edge not used in T
- ▶ Let C be the cycle created by adding uv to T
- ▶ Then uv is the **most expensive edge** of C
(or tied for most expensive, if there are ties)

Why?

- ▶ Suppose not. Let uv be an edge not used in T , that is not the most expensive in its cycle C in $T + uv$
- ▶ Create a tree T' by adding uv , then deleting the most expensive edge $u'v'$ from C .
- ▶ Then T' is a cheaper spanning tree than T , a contradiction.

Min-cost spanning trees

Min-cost spanning trees: Cycle rule

Let $G = (V, E)$ be a graph with edge weights, T a spanning tree in G . Then T is **min-cost** if and only if the following holds:

- ▶ Let $uv \in E$ be **any** edge not used in T
- ▶ Let C be the cycle created by adding uv to T
- ▶ Then uv is the **most expensive edge** of C
(or tied for most expensive, if there are ties)

Why?

- ▶ Suppose not. Let uv be an edge not used in T , that is not the most expensive in its cycle C in $T + uv$
- ▶ Create a tree T' by adding uv , then deleting the most expensive edge $u'v'$ from C .
- ▶ Then T' is a cheaper spanning tree than T , a contradiction.

Min-cost spanning trees

Min-cost spanning trees: Cycle rule

Let $G = (V, E)$ be a graph with edge weights, T a spanning tree in G . Then T is **min-cost** if and only if the following holds:

- ▶ Let $uv \in E$ be **any** edge not used in T
- ▶ Let C be the cycle created by adding uv to T
- ▶ Then uv is the **most expensive edge** of C (or tied for most expensive, if there are ties)

Why? ¹

- ▶ Suppose not. Let uv be an edge not used in T , that is not the most expensive in its cycle C in $T + uv$
- ▶ Create a tree T' by adding uv , then deleting the most expensive edge $u'v'$ from C .
- ▶ Then T' is a cheaper spanning tree than T , a contradiction.

¹This was **necessity**. To show **sufficiency** is a little bit harder.

Min-cost spanning trees: Option II

Min-cost spanning trees: Cut rule

Let $G = (V, E)$ be a graph with edge weights, T a spanning tree in G . Then T is **min-cost** if and only if the following holds:

- ▶ For **any** split of V into two parts $V = A \cup B$,
 T contains a **cheapest edge** between A and B .

Proof: Same as cycle rule (if the statement were false, you could create a cheaper tree from T as $T - uv + u'v'$).

Min-cost spanning trees: Option II

Min-cost spanning trees: Cut rule

Let $G = (V, E)$ be a graph with edge weights, T a spanning tree in G . Then T is **min-cost** if and only if the following holds:

- ▶ For **any** split of V into two parts $V = A \cup B$,
 T contains a **cheapest edge** between A and B .

Proof: Same as cycle rule (if the statement were false, you could create a cheaper tree from T as $T - uv + u'v'$).

Minimum spanning tree: algorithms

- ▶ Prim's algorithm (the one you know):
 1. Begin with tree T of single vertex v , no edges
 2. Repeat until T is a spanning tree:
 - 2.1 Extend T by the **cheapest** edge uv where u in T , v not in T
- ▶ Kruskal's algorithm
 1. Begin with $T = \emptyset$
 2. Sort all edges of E : $w(e_1) \leq w(e_2) \leq \dots \leq w(e_m)$
 3. For every edge e in this order:
 - 3.1 If adding e to T does not create a cycle, add e to T .

Minimum spanning tree: algorithms


- ▶ Prim's algorithm (the one you know):
 1. Begin with tree T of single vertex v , no edges
 2. Repeat until T is a spanning tree:
 - 2.1 Extend T by the **cheapest** edge uv where u in T , v not in T
- ▶ Kruskal's algorithm
 1. Begin with $T = \emptyset$
 2. Sort all edges of E : $w(e_1) \leq w(e_2) \leq \dots \leq w(e_m)$
 3. For every edge e in this order:
 - 3.1 If adding e to T does not create a cycle, add e to T .

Minimum spanning tree: algorithms


- ▶ Prim's algorithm (the one you know):
 1. Begin with tree T of single vertex v , no edges
 2. Repeat until T is a spanning tree:
 - 2.1 Extend T by the **cheapest** edge uv where u in T , v not in T
- ▶ Kruskal's algorithm
 1. Begin with $T = \emptyset$
 2. Sort all edges of E : $w(e_1) \leq w(e_2) \leq \dots \leq w(e_m)$
 3. For every edge e in this order:
 - 3.1 If adding e to T does not create a cycle, add e to T .
- ▶ Issues:
 - ▶ Algorithm workings
 - ▶ Correctness
 - ▶ Efficient implementation

Using the algorithms


(1) Run Prim's algorithm (start from vertex a) and Kruskal's algorithm on graph G below. (2) Using Kruskal's algorithm, determine how many minimum spanning trees (MSTs) graph F has. (3) How many MSTs are in the graph H' obtained from H by replacing cost 11 with 1 and 12 with 2?



G



H



F

Using the algorithms: Solutions

1. Graph G has a minimum spanning tree (MST) with the edges of costs 1, 2, 3, 4, 5.
2. Since in any MSP of F only 2 edges of the 3 edges of cost 1 can be used and the same for edges of cost 6, F has $3 \times 3 = 9$ MSTs (edges of cost 3 and 4 must be in each MST).
3. In H each MST must have edges of cost 1, but only one of 3 edges of cost 2. Hence, 3 MSTs.

Using the algorithms: Solutions

1. Graph G has a minimum spanning tree (MST) with the edges of costs 1, 2, 3, 4, 5.
2. Since in any MSP of F only 2 edges of the 3 edges of cost 1 can be used and the same for edges of cost 6, F has $3 \times 3 = 9$ MSTs (edges of cost 3 and 4 must be in each MST).
3. In H each MST must have edges of cost 1, but only one of 3 edges of cost 2. Hence, 3 MSTs.

Using the algorithms: Solutions

1. Graph G has a minimum spanning tree (MST) with the edges of costs 1, 2, 3, 4, 5.
2. Since in any MSP of F only 2 edges of the 3 edges of cost 1 can be used and the same for edges of cost 6, F has $3 \times 3 = 9$ MSTs (edges of cost 3 and 4 must be in each MST).
3. In H each MST must have edges of cost 1, but only one of 3 edges of cost 2. Hence, 3 MSTs.

Spanning trees: Prim's algorithm

Prim's algorithm, outline

Input: Graph $G = (V, E)$, edge weights w

- ▶ Start with a tree T of a single vertex v , no edges
- ▶ Repeat until complete:
 - ▶ Extend T with the **cheapest** edge e adding a new vertex to T
- ▶ Discover/Extend central loop (Priority Queue)

Prim's algorithm, outline

Input: Graph $G = (V, E)$, edge weights w

- ▶ Start with a tree T of a single vertex v , no edges
- ▶ Repeat until complete:
 - ▶ Extend T with the **cheapest** edge e adding a new vertex to T
- ▶ **Discover/Extend** central loop (Priority Queue)

First implementation

- ▶ Data structures:
 1. Vertices **marked/unmarked**, initially all unmarked
 2. **Priority Queue pq of edges**, initially empty
 3. The tree **T**
- ▶ Use a “visit” subroutine **process(v)**:
 1. Mark v as visited
 2. For every unmarked neighbour u of v :
 - 2.1 Add edge uv to **pq** with weight $w(uv)$
- ▶ Code, main loop:
 1. Select initial vertex v , **process(v)**
 2. Until **pq** is empty:
 - 2.1 **$e = \text{pq.deleteMin}()$** // pull cheapest edge e from **pq**
 - 2.2 If e has an unmarked endpoint u :
add e to T , **process(u)**

First implementation

- ▶ Data structures:
 1. Vertices **marked/unmarked**, initially all unmarked
 2. **Priority Queue pq of edges**, initially empty
 3. The tree **T**
- ▶ Use a “visit” subroutine **process(v)**:
 1. Mark v as visited
 2. For every unmarked neighbour u of v :
 - 2.1 Add edge uv to **pq** with weight $w(uv)$
- ▶ Code, main loop:
 1. Select initial vertex v , **process(v)**
 2. Until **pq** is empty:
 - 2.1 **$e = \text{pq.deleteMin}()$** // pull cheapest edge e from **pq**
 - 2.2 If e has an unmarked endpoint u :
add e to T , **process(u)**

First implementation

- ▶ Data structures:
 1. Vertices **marked/unmarked**, initially all unmarked
 2. **Priority Queue pq of edges**, initially empty
 3. The tree **T**
- ▶ Use a “visit” subroutine **process(v)**:
 1. Mark v as visited
 2. For every unmarked neighbour u of v :
 - 2.1 Add edge uv to **pq** with weight $w(uv)$
- ▶ Code, main loop:
 1. Select initial vertex v , **process(v)**
 2. Until **pq** is empty:
 - 2.1 **$e = \text{pq.deleteMin}()$** // pull cheapest edge e from **pq**
 - 2.2 If e has an unmarked endpoint u :
add e to T , **process(u)**

First version, running time

- ▶ Our code performs:
 1. $\mathcal{O}(|V|)$ executions of “visit all neighbours of a vertex”, takes $\mathcal{O}(|E|)$ time in total
 2. $\mathcal{O}(|E|)$ insertions into **pq** at $\mathcal{O}(\log |E|)$ per op
 3. $\mathcal{O}(|E|)$ deleteMin() calls to **pq** at $\mathcal{O}(\log |E|)$ per op
- ▶ Total time² $\mathcal{O}(|E| \log |E|)$
- ▶ Correctness proof (skipped):
 - ▶ Can show that the tree constructed follows **cut rule**

²We assume that the graph is connected, otherwise it has no spanning tree

– so $|V| \leq |E| - 1$

First version, running time

- ▶ Our code performs:
 1. $\mathcal{O}(|V|)$ executions of “visit all neighbours of a vertex”, takes $\mathcal{O}(|E|)$ time in total
 2. $\mathcal{O}(|E|)$ insertions into **pq** at $\mathcal{O}(\log |E|)$ per op
 3. $\mathcal{O}(|E|)$ deleteMin() calls to **pq** at $\mathcal{O}(\log |E|)$ per op
- ▶ Total time² $\mathcal{O}(|E| \log |E|)$
- ▶ Correctness proof (skipped):
 - ▶ Can show that the tree constructed follows **cut rule**

²We assume that the graph is connected, otherwise it has no spanning tree

– so $|V| \leq |E| - 1$

Advanced implementation

- ▶ The code adds **edges** to pq, but the algorithm needs to discover **vertices**
- ▶ This leads to a minor inefficiency (in space and time)
- ▶ Using an **advanced priority queue** (with decreaseKey method) we can write the code using a **priority queue of vertices** for space $\mathcal{O}(|V|)$, time $\mathcal{O}(|E| + |V| \log |V|)$ (theoretical)
- ▶ Details at the end of slide set (optional, not examined, i.e. marked ★)

Advanced implementation

- ▶ The code adds **edges** to pq, but the algorithm needs to discover **vertices**
- ▶ This leads to a minor inefficiency (in space and time)
- ▶ Using an **advanced priority queue** (with decreaseKey method) we can write the code using a **priority queue of vertices** for space $\mathcal{O}(|V|)$, time $\mathcal{O}(|E| + |V| \log |V|)$ (theoretical)
- ▶ Details at the end of slide set (optional, not examined, i.e. marked ★)

Advanced implementation

- ▶ The code adds **edges** to pq, but the algorithm needs to discover **vertices**
- ▶ This leads to a minor inefficiency (in space and time)
- ▶ Using an **advanced priority queue** (with decreaseKey method) we can write the code using a **priority queue of vertices** for space $\mathcal{O}(|V|)$, time $\mathcal{O}(|E| + |V| \log |V|)$ (theoretical)
- ▶ Details at the end of slide set (optional, not examined, i.e. marked ★)

Kruskal's algorithm

Kruskal's algorithm

The pseudocode looks even simpler than Prim:

1. Begin with $T = \emptyset$
2. Sort all edges of E : $w(e_1) \leq w(e_2) \leq \dots \leq w(e_m)$
3. For every edge e in this order:
 - 3.1 If adding e to T does not create a cycle, add e to T .

Note: During the loop, T has several connected components – it's a forest, not a tree

Kruskal's algorithm

The pseudocode looks even simpler than Prim:

1. Begin with $T = \emptyset$
2. Sort all edges of E : $w(e_1) \leq w(e_2) \leq \dots \leq w(e_m)$
3. For every edge e in this order:
 - 3.1 If adding e to T does not create a cycle,
add e to T .

Note: During the loop, T has **several connected components** – it's a **forest**, not a tree

Correctness

Kruskal's algorithm creates a minimum spanning tree

- ▶ It clearly creates a spanning tree
- ▶ Min-cost via the **cycle rule**:
 - ▶ Let e be any edge not used in T
 - ▶ Since e was not added to T , at the point where it was considered by the algorithm, it would have created a cycle C
 - ▶ Since the edges are considered in sorted order, e was a most-expensive edge of C
- ▶ Every edge e from G not in T satisfies the cycle rule \Rightarrow
 T is a minimum spanning tree

Correctness

Kruskal's algorithm creates a minimum spanning tree

- ▶ It clearly creates a spanning tree
- ▶ Min-cost via the **cycle rule**:
 - ▶ Let e be any edge not used in T
 - ▶ Since e was not added to T , at the point where it was considered by the algorithm, it would have created a cycle C
 - ▶ Since the edges are considered in sorted order, e was a most-expensive edge of C
- ▶ Every edge e from G not in T satisfies the cycle rule \Rightarrow
 T is a minimum spanning tree

Correctness

Kruskal's algorithm creates a minimum spanning tree

- ▶ It clearly creates a spanning tree
- ▶ Min-cost via the **cycle rule**:
 - ▶ Let e be any edge not used in T
 - ▶ Since e was not added to T , at the point where it was considered by the algorithm, it would have created a cycle C
 - ▶ Since the edges are considered in sorted order, e was a most-expensive edge of C
- ▶ Every edge e from G not in T satisfies the cycle rule \Rightarrow
 T is a minimum spanning tree

Correctness

Kruskal's algorithm creates a minimum spanning tree

- ▶ It clearly creates a spanning tree
- ▶ Min-cost via the **cycle rule**:
 - ▶ Let e be any edge not used in T
 - ▶ Since e was not added to T , at the point where it was considered by the algorithm, it would have created a cycle C
 - ▶ Since the edges are considered in sorted order, e was a most-expensive edge of C
- ▶ Every edge e from G not in T satisfies the cycle rule \Rightarrow
 T is a minimum spanning tree

Correctness

Kruskal's algorithm creates a minimum spanning tree

- ▶ It clearly creates a spanning tree
- ▶ Min-cost via the **cycle rule**:
 - ▶ Let e be any edge not used in T
 - ▶ Since e was not added to T , at the point where it was considered by the algorithm, it would have created a cycle C
 - ▶ Since the edges are considered in sorted order, e was a most-expensive edge of C
- ▶ Every edge e from G not in T satisfies the cycle rule \Rightarrow
 T is a minimum spanning tree

Correctness

Kruskal's algorithm creates a minimum spanning tree

- ▶ It clearly creates a spanning tree
- ▶ Min-cost via the **cycle rule**:
 - ▶ Let e be any edge not used in T
 - ▶ Since e was not added to T , at the point where it was considered by the algorithm, it would have created a cycle C
 - ▶ Since the edges are considered in sorted order, e was a most-expensive edge of C
- ▶ Every edge e from G not in T satisfies the cycle rule \Rightarrow
 T is a minimum spanning tree

Implementation

1. Sort the edges
 - ▶ No problem
2. For every edge uv in sorted order:
 - 2.1 Decide if it completes a cycle
 - 2.2 If it does not, add it to T

Completes a cycle has two implementations:

- ▶ A slow one:
 - ▶ Search through the graph formed by T , starting from u , to see if we find v (e.g., DFS from u in T)
 - ▶ Up to $\mathcal{O}(n)$ time for a single check
- ▶ A fast one: Using a data structure called Union Find

Implementation

1. Sort the edges
 - ▶ No problem
2. For every edge uv in sorted order:
 - 2.1 Decide if it completes a cycle
 - 2.2 If it does not, add it to T

Completes a cycle has two implementations:

- ▶ A slow one:
 - ▶ Search through the graph formed by T , starting from u , to see if we find v (e.g., DFS from u in T)
 - ▶ Up to $\mathcal{O}(n)$ time for a single check
- ▶ A fast one: Using a data structure called Union Find

Implementation

1. Sort the edges
 - ▶ No problem
2. For every edge uv in sorted order:
 - 2.1 Decide if it completes a cycle
 - 2.2 If it does not, add it to T

Completes a cycle has two implementations:

- ▶ A slow one:
 - ▶ Search through the graph formed by T , starting from u , to see if we find v (e.g., DFS from u in T)
 - ▶ Up to $\mathcal{O}(n)$ time for a single check
- ▶ A fast one: Using a data structure called Union Find

Fast implementation

We need a data structure for connected components:

- ▶ **Find(v)** – return a name for the connected component of T that contains v
- ▶ **Union(u, v)** – “merge” the components of u and v into one (by adding uv to T)

Efficient implementation using this:

1. Sort the edges E (e.g., merge sort) at $\mathcal{O}(|E| \log |E|)$ time
2. Initiate your **Union-Find** structure to “empty”
3. For every edge uv in sorted order:
 - 3.1 If **Find(u)** == **Find(v)**: ignore edge
 - 3.2 Otherwise call **Union(u, v)** and add to T

With a good **Union-Find**, the time is dominated by $\mathcal{O}(|E| \log |E|)$.

Fast implementation

We need a data structure for connected components:

- ▶ **Find**(v) – return a name for the connected component of T that contains v
- ▶ **Union**(u, v) – “merge” the components of u and v into one (by adding uv to T)

Efficient implementation using this:

1. Sort the edges E (e.g., merge sort) at $\mathcal{O}(|E| \log |E|)$ time
2. Initiate your **Union-Find** structure to “empty”
3. For every edge uv in sorted order:
 - 3.1 If **Find**(u) == **Find**(v): ignore edge
 - 3.2 Otherwise call **Union**(u, v) and add to T

With a good **Union-Find**, the time is dominated by $\mathcal{O}(|E| \log |E|)$.

Fast implementation

We need a data structure for connected components:

- ▶ **Find**(v) – return a name for the connected component of T that contains v
- ▶ **Union**(u, v) – “merge” the components of u and v into one (by adding uv to T)

Efficient implementation using this:

1. Sort the edges E (e.g., merge sort) at $\mathcal{O}(|E| \log |E|)$ time
2. Initiate your **Union-Find** structure to “empty”
3. For every edge uv in sorted order:
 - 3.1 If **Find**(u) == **Find**(v): ignore edge
 - 3.2 Otherwise call **Union**(u, v) and add to T

With a good **Union-Find**, the time is dominated by $\mathcal{O}(|E| \log |E|)$.

★ Prim: advanced implementation

“Advanced PQ” interface

- ▶ In addition to the usual min-Priority Queue operations:
 - ▶ `insert(item, weight)`: add `item` with value `weight`
 - ▶ `minItem()`, `deleteMin()`: Return min-weight item
- ▶ We will need a Priority Queue that also supports the following:
 1. `contains(item)`: membership test
 2. `currentValue(item)`: return current item priority/weight
 3. `decreaseValue(item, newWeight)`: decrease weight of `item` to `newWeight`
- ▶ Implementations:
 1. `Heap+Hash table` with $\mathcal{O}(\log n)$ time per `decreaseKey`
 2. `Fibonacci heaps`: Advanced structure, not yet practical
 - ▶ Impractical: Pointer jumping, complex code
 - ▶ But average time $\mathcal{O}(1)$ for `decreaseKey`

“Advanced PQ” interface

- ▶ In addition to the usual min-Priority Queue operations:
 - ▶ `insert(item, weight)`: add `item` with value `weight`
 - ▶ `minItem()`, `deleteMin()`: Return min-weight item
- ▶ We will need a Priority Queue that also supports the following:
 1. `contains(item)`: membership test
 2. `currentValue(item)`: return current item priority/weight
 3. `decreaseValue(item, newWeight)`: decrease weight of `item` to `newWeight`
- ▶ Implementations:
 1. `Heap+Hash table` with $\mathcal{O}(\log n)$ time per `decreaseKey`
 2. `Fibonacci heaps`: Advanced structure, not yet practical
 - ▶ Impractical: Pointer jumping, complex code
 - ▶ But average time $\mathcal{O}(1)$ for `decreaseKey`

“Advanced PQ” interface

- ▶ In addition to the usual min-Priority Queue operations:
 - ▶ `insert(item, weight)`: add `item` with value `weight`
 - ▶ `minItem()`, `deleteMin()`: Return min-weight item
- ▶ We will need a Priority Queue that also supports the following:
 1. `contains(item)`: membership test
 2. `currentValue(item)`: return current item priority/weight
 3. `decreaseValue(item, newWeight)`: decrease weight of `item` to `newWeight`
- ▶ Implementations:
 1. `Heap+Hash table` with $\mathcal{O}(\log n)$ time per `decreaseKey`
 2. `Fibonacci heaps`: Advanced structure, not yet practical
 - ▶ Impractical: Pointer jumping, complex code
 - ▶ But average time $\mathcal{O}(1)$ for `decreaseKey`

“Advanced PQ” interface

- ▶ In addition to the usual min-Priority Queue operations:
 - ▶ `insert(item, weight)`: add `item` with value `weight`
 - ▶ `minItem()`, `deleteMin()`: Return min-weight item
- ▶ We will need a Priority Queue that also supports the following:
 1. `contains(item)`: membership test
 2. `currentValue(item)`: return current item priority/weight
 3. `decreaseValue(item, newWeight)`: decrease weight of `item` to `newWeight`
- ▶ Implementations:
 1. `Heap+Hash table` with $\mathcal{O}(\log n)$ time per `decreaseKey`
 2. `Fibonacci heaps`: Advanced structure, not yet practical
 - ▶ Impractical: Pointer jumping, complex code
 - ▶ But average time $\mathcal{O}(1)$ for `decreaseKey`

“Advanced PQ” interface

- ▶ In addition to the usual min-Priority Queue operations:
 - ▶ `insert(item, weight)`: add `item` with value `weight`
 - ▶ `minItem()`, `deleteMin()`: Return min-weight item
- ▶ We will need a Priority Queue that also supports the following:
 1. `contains(item)`: membership test
 2. `currentValue(item)`: return current item priority/weight
 3. `decreaseValue(item, newWeight)`: decrease weight of `item` to `newWeight`
- ▶ Implementations:
 1. `Heap+Hash table` with $\mathcal{O}(\log n)$ time per `decreaseKey`
 2. `Fibonacci heaps`: Advanced structure, not yet practical
 - ▶ Impractical: Pointer jumping, complex code
 - ▶ But average time $\mathcal{O}(1)$ for `decreaseKey`

Advanced Prim implementation: Outline

- ▶ Data structures:

1. Min-Priority Queue pq **containing vertices**
2. **edgeTo[v]** array codes edges, will encode tree
3. Vertices marked/unmarked as before

- ▶ Principles:

- ▶ pq contains candidate vertex to add to tree
- ▶ **Priority** (weight) of vertex v in pq is **cheapest known edge** to v from tree
- ▶ When we pull v from pq (“discover” step), we **update** neighbours u of v if a **cheaper** edge to u was found
- ▶ **If** a cheaper edge was found, we call **decreaseKey** to v instead of a full **insert**

Advanced Prim implementation: Outline

- ▶ Data structures:

1. Min-Priority Queue pq containing vertices
2. `edgeTo[v]` array codes edges, will encode tree
3. Vertices marked/unmarked as before

- ▶ Principles:

- ▶ pq contains candidate vertex to add to tree
- ▶ Priority (weight) of vertex v in pq is cheapest known edge to v from tree
- ▶ When we pull v from pq (“discover” step), we update neighbours u of v if a cheaper edge to u was found
- ▶ If a cheaper edge was found, we call `decreaseKey` to v instead of a full `insert`

Advanced Prim implementation: Outline

- ▶ Data structures:

1. Min-Priority Queue pq containing vertices
2. `edgeTo[v]` array codes edges, will encode tree
3. Vertices marked/unmarked as before

- ▶ Principles:

- ▶ pq contains candidate vertex to add to tree
- ▶ Priority (weight) of vertex v in pq is cheapest known edge to v from tree
- ▶ When we pull v from pq (“discover” step), we update neighbours u of v if a cheaper edge to u was found
- ▶ If a cheaper edge was found, we call `decreaseKey` to v instead of a full `insert`

Full code

- ▶ Data structures:
 1. Min-Priority Queue `pq` containing vertices
 2. `edgeTo[v]` array containing edges, will encode tree
- ▶ Main loop:
 1. Select initial vertex v , call `pq.insert(v,0)`
 2. Until `pq` is empty:
 - 2.1 `u=pq.deleteMin()`; mark u
 - 2.2 For every neighbour v of u :
Call `Update(v, w(uv))`
- ▶ Support code: `Update(Vertex v, Edge uv)`:
 1. If v marked: return, do nothing
 2. If v not in `pq`:
 - 2.1 `edgeTo[v] = uv`
 - 2.2 `pq.insert(v, w(uv))`
 3. If `w(uv) < pq.currentValue(v)`:
 - 3.1 `edgeTo[v] = uv`
 - 3.2 `pq.decreaseValue(v, w(uv))`

Full code

- ▶ Data structures:
 1. Min-Priority Queue `pq` containing vertices
 2. `edgeTo[v]` array containing edges, will encode tree
- ▶ Main loop:
 1. Select initial vertex v , call `pq.insert(v,0)`
 2. Until `pq` is empty:
 - 2.1 `u=pq.deleteMin()`; mark u
 - 2.2 For every neighbour v of u :
Call `Update(v, w(uv))`
- ▶ Support code: `Update(Vertex v, Edge uv)`:
 1. If v marked: return, do nothing
 2. If v not in `pq`:
 - 2.1 `edgeTo[v] = uv`
 - 2.2 `pq.insert(v, w(uv))`
 3. If `w(uv) < pq.currentValue(v)`:
 - 3.1 `edgeTo[v] = uv`
 - 3.2 `pq.decreaseValue(v, w(uv))`

Full code

- ▶ Data structures:
 1. Min-Priority Queue `pq` containing vertices
 2. `edgeTo[v]` array containing edges, will encode tree
- ▶ Main loop:
 1. Select initial vertex v , call `pq.insert(v,0)`
 2. Until `pq` is empty:
 - 2.1 `u=pq.deleteMin()`; mark u
 - 2.2 For every neighbour v of u :
Call `Update(v, w(uv))`
- ▶ Support code: `Update(Vertex v, Edge uv)`:
 1. If v marked: return, do nothing
 2. If v not in `pq`:
 - 2.1 `edgeTo[v] = uv`
 - 2.2 `pq.insert(v, w(uv))`
 3. If `w(uv) < pq.currentValue(v)`:
 - 3.1 `edgeTo[v] = uv`
 - 3.2 `pq.decreaseValue(v, w(uv))`

Advanced version: running time

- ▶ Our new implementation performs:
 1. $\mathcal{O}(|V|)$ insertions into pq
 2. $\mathcal{O}(|E|)$ update, decreaseKey calls
 3. $\mathcal{O}(|V|)$ deletions from pq
- ▶ With our familiar implementation:
 1. All operations take $\mathcal{O}(\log |V|)$ time
 2. Worst-case $\mathcal{O}(|E| \log |V|) = \mathcal{O}(|E| \log |E|)$ as before
 3. New **best case**: Zero **decreaseKey** calls,
 $\mathcal{O}(|E| + |V| \log |V|)$ time (maybe even just $\mathcal{O}(|E|)$?)
- ▶ With Fibonacci heaps (theoretical result):
 1. Insert/delete amounts to $\mathcal{O}(|V| \log |V|)$ time
 2. $\mathcal{O}(|E|)$ times **decreaseKey** takes $\mathcal{O}(|E|)$ total time
 3. In total **worst-case time** $\mathcal{O}(|E| + |V| \log |V|)$

Advanced version: running time

- ▶ Our new implementation performs:
 1. $\mathcal{O}(|V|)$ insertions into pq
 2. $\mathcal{O}(|E|)$ update, decreaseKey calls
 3. $\mathcal{O}(|V|)$ deletions from pq
- ▶ With our familiar implementation:
 1. All operations take $\mathcal{O}(\log |V|)$ time
 2. Worst-case $\mathcal{O}(|E| \log |V|) = \mathcal{O}(|E| \log |E|)$ as before
 3. New **best case**: Zero **decreaseKey** calls,
 $\mathcal{O}(|E| + |V| \log |V|)$ time (maybe even just $\mathcal{O}(|E|)$?)
- ▶ With Fibonacci heaps (theoretical result):
 1. Insert/delete amounts to $\mathcal{O}(|V| \log |V|)$ time
 2. $\mathcal{O}(|E|)$ times **decreaseKey** takes $\mathcal{O}(|E|)$ total time
 3. In total **worst-case time** $\mathcal{O}(|E| + |V| \log |V|)$

Advanced version: running time

- ▶ Our new implementation performs:
 1. $\mathcal{O}(|V|)$ insertions into pq
 2. $\mathcal{O}(|E|)$ update, decreaseKey calls
 3. $\mathcal{O}(|V|)$ deletions from pq
- ▶ With our familiar implementation:
 1. All operations take $\mathcal{O}(\log |V|)$ time
 2. Worst-case $\mathcal{O}(|E| \log |V|) = \mathcal{O}(|E| \log |E|)$ as before
 3. New **best case**: Zero **decreaseKey** calls,
 $\mathcal{O}(|E| + |V| \log |V|)$ time (maybe even just $\mathcal{O}(|E|)$?)
- ▶ With Fibonacci heaps (theoretical result):
 1. Insert/delete amounts to $\mathcal{O}(|V| \log |V|)$ time
 2. $\mathcal{O}(|E|)$ times **decreaseKey** takes $\mathcal{O}(|E|)$ total time
 3. In total **worst-case time** $\mathcal{O}(|E| + |V| \log |V|)$